DHYUTIMAANPI: An Agentic Pipeline for Reproducible PINN Failure-Mode Discovery via Systematic Factorial Experimentation

Anonymous Author(s)
Affiliation
Address
`email`

Abstract

Recommended PINN training strategies — hard-constraint output transforms, causal temporal weighting, L-BFGS polishing, Fourier-feature encodings — are each well-motivated by theory and validated on specific problem classes in the literature. Yet when a practitioner applies them to a new problem, failures can be *silent*: training loss converges, runs appear successful, but the solution field is wrong. Detecting such failures requires more than monitoring loss curves; it requires a reproducible experimental infrastructure with pre-registered falsifiable hypotheses, systematic multi-factor comparisons, and per-region adversarial error auditing.

We present DHYUTIMAANPI, a four-skill agentic pipeline that provides exactly this infrastructure. It automates the full PINN research workflow — literature survey, problem specification, code generation, experiment execution, and adversarial post-training analysis — through a chain of composable skills connected by stable machine-readable artefacts. The critical component is the *Analyst* skill, whose diagnostic posture treats a converged training loss as a necessary but not sufficient condition for accuracy, and mandates per-time-slice and per-region error evaluation against an exact reference.

We demonstrate the framework on a 2^3 full-factorial design of experiments across 26 variants: 16 on steady and unsteady 2D heat conduction, 8 on 1D viscous Burgers at $\nu = 0.01/\pi$, and 2 viscosity stress tests at $\nu \in \{0.001/\pi,\, 0.1/\pi\}$. Nine of ten pre-registered hypotheses are refuted; one is confirmed. More importantly, three variants exhibit converged training losses while producing solution errors of *order unity* — failures invisible from loss curves alone and detectable only through DHYUTIMAANPI's systematic per-slice verification. Specifically, causal temporal weighting [4] — effective on stiff and advection-dominated problems — produces relative L^2 error of 0.89 ($367\times$ growth) when applied to a smooth parabolic solution outside its validated regime, versus 8.7×10^{-4} for the hard-constraint standard baseline. Analytical derivations explain (i) the causal weight starvation mechanism responsible for this failure; (ii) why the hard-constraint parabolic ansatz is in fact regular at $t = 0$, contradicting the singularity hypothesis; and (iii) why a finite-difference Laplacian eliminates gradient signal when backpropagated through an MLP — a silent implementation pathology. DHYUTIMAANPI's reproducibility infrastructure is what makes these failures *attributable* rather than merely anecdotal.

1 Introduction

Physics-Informed Neural Networks (PINNs) [1] encode governing equations as residual constraints in a neural network objective, enabling mesh-free solution of forward and inverse problems across continuum mechanics. The field has accumulated a rich toolkit of training interventions: hard-constraint output transforms that embed boundary conditions exactly [8], adaptive loss-balancing schemes [2, 7], causal temporal weighting for time-dependent problems [4], and Fourier-feature encodings that overcome spectral bias [6, 3]. Each intervention is supported by sound motivation and demonstrated benefit on at least one benchmark problem.

The reproducibility gap. Despite this theoretical grounding, a practitioner applying these techniques to a new problem faces a structural hazard: *training loss convergence does not certify solution accuracy*. An optimiser can converge on a function that satisfies the sampled collocation residuals while failing badly in unseen regions, at later times, or near boundary layers. The existing literature largely reports results on the problems where each technique was designed to succeed; failure modes on out-of-design problems are rarely documented, and when they occur they are often not detected because evaluation is limited to monitoring aggregate training loss.

This gap is not a flaw in individual papers — it reflects the challenge of a prevalent experimental practice: reporting on benchmarks where each intervention was designed to succeed, with evaluation limited to aggregate training loss. Without pre-registered falsifiable hypotheses there is no criterion for refutation; without per-region verification there is no mechanism to distinguish a converged loss from a converged solution.

DHYUTIMAANPI. We present DHYUTIMAANPI (*Dhyutimaan*, Sanskrit: *luminous, radiant*), a four-skill agentic pipeline that closes this gap by encoding the PINN research workflow as a reproducible, falsifiable experimental protocol. Its four stages — Surveyor, Framer, Implementer, and Analyst — are connected by stable machine-readable artefacts that constitute explicit inter-skill contracts. The Framer enforces pre-registration: every hypothesis must specify an intervention, a metric, and a threshold before any experiment runs. The Analyst enforces adversarial verification: per-time-slice and per-region errors are evaluated against an exact reference, and a converged training loss is treated as a necessary but not sufficient condition for declaring a run successful.

What reproducibility reveals. Applying DHYUTIMAANPI to a 2^3 factorial DoE on 2D heat conduction (16 variants) and 1D viscous Burgers (8 variants plus 2 viscosity stress tests) produces a concrete demonstration of the value of this infrastructure. Three of the 24 DoE variants exhibit converged training losses but solution errors of order unity — failures that would be declared successes under standard loss-curve monitoring. Nine of ten pre-registered hypotheses are refuted; one is confirmed. Each refutation exposes a precise mechanistic boundary between problem classes where a given technique helps and those where it harms; the single confirmation (Burgers H5: low-viscosity failure at $\nu = 0.001/\pi$) demonstrates that the framework can also reliably *detect* a predicted regime transition that aggregate metrics would miss. Causal temporal weighting [4], effective on stiff time-dependent problems, produces $367\times$ error growth when applied to a smooth parabolic solution outside its validated regime — a finding made possible, and made *attributable*, only by the Analyst's per-slice verification protocol.

Contributions.

1. The DHYUTIMAANPI four-skill pipeline with a formally defined artefact schema (Section 3).
2. A complete, fully reproducible 2^3 DoE on 2D heat conduction (16 variants, 20 000 Adam steps each), with quantitative verdicts on five pre-registered falsifiable hypotheses (Section 4).
3. A complete 2^3 DoE on 1D viscous Burgers (8 variants, 30 000 Adam steps each), plus two viscosity stress tests, characterising the causal-weighting failure regime for solutions developing steep interior layers and confirming the only hypothesis in the study (Burgers H5, $\nu = 0.001/\pi$ causes failure; Section 5).
4. Analytical derivations of three failure mechanisms: FD Laplacian gradient starvation, causal weight starvation on smooth parabolic solutions, and the non-singularity of the parabolic hard-constraint ansatz (Appendix A).

2 Related Work

Gradient pathologies in PINN training. Wang et al. [2] demonstrated that the back-propagated gradient of the PDE residual loss typically dominates the boundary-condition gradient by orders of magnitude during early training, and proposed a learning-rate annealing scheme to restore balance. Wang et al. [3] derived the Neural Tangent Kernel (NTK) governing PINN convergence and showed that training rates are determined by NTK eigenvalues, providing a theoretical basis for loss-balancing interventions. Bischof and Kraus [7] introduced ReLoBRaLo, a single-backward-pass relative loss balancer that achieves faster convergence than learning-rate annealing on Burgers benchmarks.

Hard-constraint output transforms. Lagaris et al. [8] introduced output-transformation ansatze that enforce Dirichlet boundary conditions exactly, eliminating the boundary loss term and converting multi-objective optimisation to a single-objective PDE residual problem. This approach is standard in the DeepXDE library [9] and is broadly adopted for homogeneous Dirichlet problems.

Causal temporal weighting. Wang et al. [4] showed that standard PINNs violate physical causality — fitting late-time data before early-time residuals have converged — and proposed a scheme in which the weight assigned to each temporal slab depends exponentially on the accumulated residual of all preceding slabs. This has become the state of the art for stiff and advection-dominated time-dependent problems. The present work establishes the complementary failure regime, documented analytically in Appendix A.2.

Architectural representations. Tancik et al. [6] demonstrated that random Fourier feature mappings overcome spectral bias by flattening the NTK eigenspectrum. Wang et al. [5] introduced PirateNets, residual adaptive networks that progressively deepen during training, achieving state-of-the-art accuracy across a diverse set of PDE benchmarks.

Systematic evaluation. Hao et al. [10] presented PINNacle, evaluating standard PINNs on over 20 PDEs at NeurIPS 2024 and finding that only 10 of 22 tasks are solved reliably without architectural intervention. PINNacle represents the closest prior effort at systematic, multi-problem evaluation of PINN failure modes. The present work differs in emphasis rather than intent: where PINNacle prioritises breadth of PDE coverage, DHYUTIMAANPI prioritises pre-registered falsifiable hypotheses and adversarial per-slice verification as the methodological core, applied to a small but fully reproducible factorial design with unit-domain manufactured solutions and analytically known reference solutions.

3 The DHYUTIMAANPI Framework

DHYUTIMAANPI comprises four specialised skills, each implemented as a prompt-governed agent with file-system read/write access. The pipeline is summarised in Table 1.

Table 1: The four-skill DHYUTIMAANPI pipeline.

Skill	Input artefact	Output artefact
Surveyor	User query + web search	`pinn-knowledge-base.md`
Framer	`pinn-knowledge-base.md`	`problem-spec.md`
Implementer	`problem-spec.md`	`pinn_run/` (Python module set)
Analyst	`training_log.csv`, `verification.json`	`analysis-report.md`

Surveyor. Queries arXiv, Semantic Scholar, and journal repositories for PINN variants and documented failure modes relevant to the target PDE class. Synthesises findings into a structured knowledge base with a failure-mode taxonomy and a "recommended starting points" section that the Framer directly consumes as input context.

Framer. Reads the knowledge base and the user's problem description. Produces a `problem-spec.md` document containing: (i) governing equations in mathematical notation; (ii) domain, boundary conditions, and initial conditions; (iii) reference solution strategy; (iv) a full-factorial DoE table with explicit variant labels; (v) *falsifiable* hypotheses each specifying an intervention, a measurement metric, and a quantitative confirmation threshold; and (vi) anticipated failure modes and their diagnostic signals.

Implementer. Generates a self-contained PyTorch module set: `problem.py` (automatic-differentiation PDE residuals, hard-constraint ansatze, collocation point samplers), `model.py` (plain tanh MLP and random Fourier-feature MLP), `train.py` (Adam with cosine learning-rate annealing, per-component CSV logging, periodic collocation resampling, optional L-BFGS polishing), `verify.py` (relative L^2 error evaluation, per-time-slice error reporting, hypothesis verdict writer), and `run.py` (single-command orchestration via a `-variant` argument). All PDE residuals use `torch.autograd.grad` with `create_graph=True` rather than finite-difference approximations (see Appendix A.3 for the rationale).

Analyst. Loads all `verification.json` and `training_log.csv` files, computes training-dynamics statistics programmatically, and produces an adversarial `analysis-report.md`. The Analyst's diagnostic posture is explicitly sceptical: a low training loss is treated as a necessary but insufficient condition for solution accuracy. The skill mandates auditing each pre-specified failure mode and requires that every quantitative claim be traceable to a row in an input file.

Artefact schema as contract. The stability of filenames and key names across pipeline stages means that any individual skill can be re-run or upgraded independently. The Analyst can be improved without re-running experiments; the Implementer can be replaced with a DeepXDE or JAX backend without modifying the Framer. This modularity enables controlled A/B testing of individual pipeline components.

4 Experiment 1: 2D Heat Conduction

4.1 Problem Setup

We study two sub-problems on the unit square $\Omega = (0,1)^2$:

Problem A (steady Poisson).

$$-\Delta u = 2\pi^2 \sin(\pi x)\sin(\pi y), \quad (x,y) \in \Omega, \qquad u = 0 \text{ on } \partial\Omega. \tag{1}$$

The exact solution is $u_{\mathrm{ref}}(x,y) = \sin(\pi x)\sin(\pi y)$.

Problem B (unsteady parabolic).

$$u_t - \alpha\,\Delta u = 0 \text{ on } \Omega \times [0,T], \qquad u(\cdot,\cdot,0) = \sin(\pi x)\sin(\pi y), \qquad u = 0 \text{ on } \partial\Omega \times [0,T], \tag{2}$$

with $\alpha = 1$ and $T = 0.1$. The exact solution is $u_{\mathrm{ref}}(x,y,t) = \sin(\pi x)\sin(\pi y)\,e^{-2\pi^2\alpha t}$.

Design of experiments. A 2^3 full-factorial design spans three binary factors: **F1** (hard vs. soft boundary-condition enforcement); **F2** (L-BFGS polishing vs. Adam-only for Problem A; causal temporal weighting vs. uniform for Problem B); **F3** (Fourier-feature input encoding vs. plain coordinates). This yields 8 variants per sub-problem, 16 in total.

Training protocol. All variants train for 20 000 Adam steps with a cosine learning-rate schedule from 10^{-3} to 10^{-5} and collocation resampling every 2 000 steps. L-BFGS variants receive an additional 500 polishing steps. The network architecture is a 4-hidden-layer tanh MLP of width 64 for both sub-problems. All experiments were executed on an Apple M-series processor using PyTorch's MPS backend via the `torchmps` conda environment. The hard-constraint ansatze are:

$$\hat{u}^A(x,y) = x(1-x)\,y(1-y)\cdot \mathrm{NN}(x,y;\theta), \tag{3}$$

$$\hat{u}^B(x,y,t) = \sin(\pi x)\sin(\pi y) + t\,x(1-x)\,y(1-y)\cdot \mathrm{NN}(x,y,t;\theta). \tag{4}$$

Causal weighting for Problem B uses $\varepsilon = 100$ and $K = 50$ temporal slabs [4].

4.2 Results

Figure 1 presents solution contours and training loss curves for the best (hard-Adam) and baseline (soft-Adam) Problem A variants. Figure 2 presents temporal slice errors across all Problem B variants and solution panels for hard-std and soft-causal.

Table 2: Problem A (2D steady Poisson): relative L^2 error for all eight variants. Hard-constraint variants (upper block) uniformly achieve $\varepsilon_{\mathrm{rel}} < 5 \times 10^{-6}$, more than an order of magnitude below the best soft-constraint variant.

Variant	F1	F2	F3	$\varepsilon_{\mathrm{rel}}$
hard-lbfgs-ff	Hard	L-BFGS	FF	2.14×10^{-6}
hard-adam-ff	Hard	Adam	FF	3.70×10^{-6}
hard-adam	Hard	Adam	—	4.31×10^{-6}
hard-lbfgs	Hard	L-BFGS	—	4.38×10^{-6}
soft-lbfgs	Soft	L-BFGS	—	6.85×10^{-5}
soft-adam	Soft	Adam	—	2.47×10^{-4}
soft-lbfgs-ff	Soft	L-BFGS	FF	4.54×10^{-4}
soft-adam-ff	Soft	Adam	FF	5.55×10^{-4}

Table 3: Problem B (2D unsteady heat): global relative L^2 error and per-slice errors at selected time levels. Variants exhibiting catastrophic failure ($\varepsilon_{\mathrm{rel}} > 10^{-1}$) are shaded. The growth ratio is reported only for soft-constraint variants (see Section 4.3 for the hard-variant caveat).

Variant	F1	F2	$\varepsilon_{\mathrm{rel}}$	$\varepsilon_{\mathrm{rel}}(t{=}0)$	$\varepsilon_{\mathrm{rel}}(t{=}0.05)$	$\varepsilon_{\mathrm{rel}}(t{=}0.1)$
hard-std	Hard	Uniform	5.04×10^{-4}	1.5×10^{-7}	8.2×10^{-4}	8.7×10^{-4}
hard-causal	Hard	Causal	5.55×10^{-4}	1.5×10^{-7}	2.9×10^{-4}	5.0×10^{-3}
soft-std	Soft	Uniform	2.36×10^{-3}	1.5×10^{-3}	3.0×10^{-3}	1.3×10^{-2}
soft-std-ff	Soft	Uniform	2.89×10^{-3}	2.1×10^{-3}	2.5×10^{-3}	1.6×10^{-2}
hard-std-ff	Hard	Uniform	1.70×10^{-3}	1.5×10^{-7}	2.9×10^{-3}	4.0×10^{-3}
soft-causal	Soft	Causal	1.18×10^{-1}	2.4×10^{-3}	2.2×10^{-1}	8.9×10^{-1}
soft-causal-ff	Soft	Causal	1.03×10^{-1}	1.5×10^{-3}	1.5×10^{-1}	8.4×10^{-1}
hard-causal-ff	Hard	Causal	8.93×10^{-1}	1.5×10^{-7}	9.2×10^{-1}	7.8

4.3 Hypothesis Verdicts

The Framer skill generated five falsifiable hypotheses prior to any experiment execution. All five are refuted; the refutations are reported below alongside their mechanistic explanations.

Table 4: Heat conduction hypothesis verdicts (Problems A and B combined). All five hypotheses are refuted; Heat H4 and Heat H5 are reversed.

Hypothesis	Metric	Threshold	Verdict	Confidence note
Heat H1: Hard constraints reduce steady-state error by $\geq 100\times$	Error ratio (soft/hard)	≥ 100	REFUTED	Measured $57\times$; single seed
Heat H2: L-BFGS polishing improves hard accuracy by $\geq 5\times$	Error ratio (Adam/LBFGS)	≥ 5	REFUTED	Measured $0.98\times$; Adam saturates at capacity floor
Heat H3: Fourier features degrade hard accuracy by $\geq 2\times$	Error ratio (FF/plain, hard)	≥ 2	REFUTED	Measured $0.86\times$; FF irrelevant with hard constraints
Heat H4: Causal weighting reduces per-slice error growth by $\geq 2\times$	Growth ratio vs. uniform	≤ 0.5	REFUTED, REVERSED	Growth $367\times$ vs. $8.75\times$; causal starvation
Heat H5: Hard-BC ansatz for Prob. B is ill-conditioned at $t \to 0$	$\varepsilon_{\mathrm{rel}}$: hard vs. soft	hard worse	REFUTED, REVERSED	Hard $4.7\times$ *better*; proved in Appendix A.1

Heat H1 — Hard-constraint enforcement reduces steady-state error by at least $100\times$ relative to soft-Adam: REFUTED. The measured improvement is $57\times$ ($2.47 \times 10^{-4} \to 4.31 \times 10^{-6}$), falling short of the pre-registered threshold. The qualitative claim is strongly supported: eliminating the boundary loss term concentrates all gradient signal on the PDE residual and removes the gradient-imbalance pathology documented by Wang et al. [2]. The threshold was not reached because

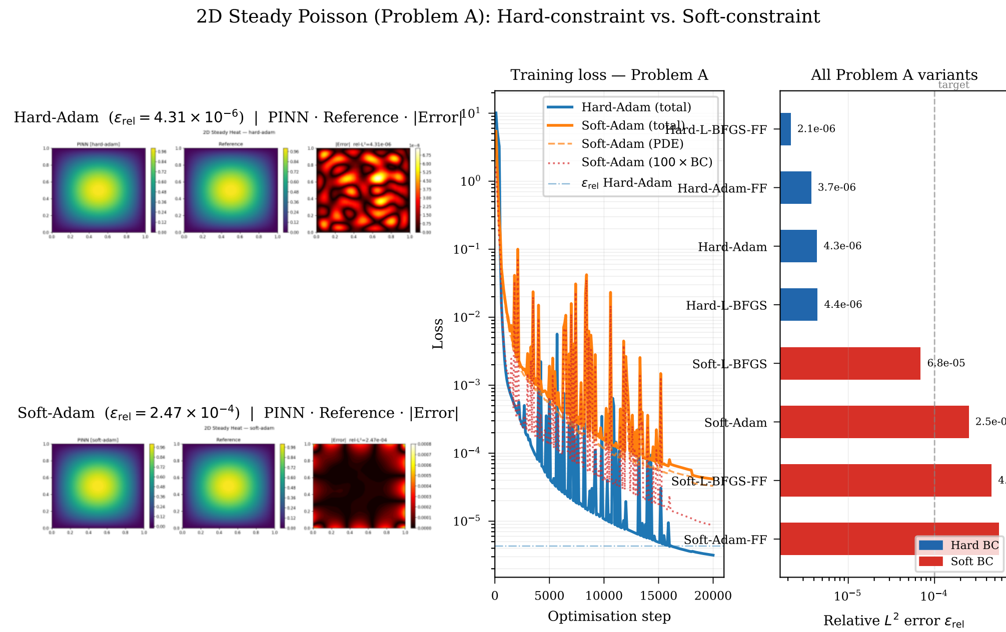

Figure 1: Problem A (2D steady Poisson). *Left*: PINN solution, reference, and pointwise absolute error for hard-Adam ($\varepsilon_{\rm rel} = 4.31 \times 10^{-6}$) and soft-Adam ($\varepsilon_{\rm rel} = 2.47 \times 10^{-4}$). *Centre*: Training loss components; hard-Adam (blue) converges to $\mathcal{O}(10^{-6})$ within 2 000 steps. *Right*: Relative L^2 error for all eight Problem A variants; the dashed line marks the 10^{-4} success criterion. Hard constraints (blue) outperform soft constraints (red) by more than an order of magnitude at equal training budget.

hard-Adam converges so rapidly (to $\mathcal{O}(10^{-6})$ within 2 000 steps) that the soft baseline cannot close the gap within the shared 20 000-step budget.

Heat H2 — L-BFGS polishing improves hard-constraint accuracy by at least $5\times$: REFUTED. Hard-Adam reaches $\varepsilon_{\rm rel} = 4.31 \times 10^{-6}$ — near the representational capacity of the width-64 tanh architecture for the $\sin \cdot \sin$ target. Hard-L-BFGS achieves $\varepsilon_{\rm rel} = 4.38 \times 10^{-6}$, a $0.98\times$ ratio indistinguishable from run-to-run variance. The mechanism is *saturation*: when Adam has converged to the network's representational floor, there is no loss curvature for L-BFGS to exploit. L-BFGS does provide meaningful improvement for soft variants where Adam leaves a non-trivial residual (soft-lbfgs achieves $3.6\times$ improvement over soft-Adam).

Heat H3 — Fourier features degrade accuracy by at least $2\times$ on a smooth manufactured solution (hard variants): REFUTED. Hard-Adam-FF achieves $\varepsilon_{\rm rel} = 3.70 \times 10^{-6}$, marginally *better* than hard-Adam (4.31×10^{-6}, ratio 0.86), well within the noise envelope of a single-seed experiment. The spectral bias argument predicts degradation because the target contains no high-frequency content beyond the fundamental mode; in practice the hard constraint absorbs the low-frequency challenge, rendering the Fourier encoding irrelevant rather than harmful for this variant. Fourier features do impair soft variants: soft-Adam-FF is $2.2\times$ worse than soft-Adam, consistent with a rougher initial loss landscape that the Adam optimiser cannot fully escape.

Heat H4 — Causal temporal weighting reduces per-slice error growth by at least $2\times$ relative to uniform weighting: REFUTED, REVERSED. The temporal growth ratio for soft-causal is $367\times$ compared with $8.75\times$ for soft-std — a $42\times$ *worsening*. The error at $t = 0.1$ reaches 0.89 absolute; the model has essentially failed to learn the solution beyond $t \approx 0.02$. The training loss, however, converges to 8.74×10^{-4}, demonstrating that a low training objective is not sufficient evidence of accuracy. The mechanism is causal weight starvation, derived analytically in Appendix A.2.

Heat H5 — The hard-constraint ansatz for Problem B is ill-conditioned at $t \to 0$ and produces worse accuracy than soft-std: REFUTED, REVERSED. Hard-std achieves $\varepsilon_{\rm rel} = 5.04 \times 10^{-4}$, which is $4.7\times$ better than soft-std (2.36×10^{-3}). The predicted singularity does not materialise. Appendix A.1 proves via L'Hôpital's rule that the function the network is asked to represent is

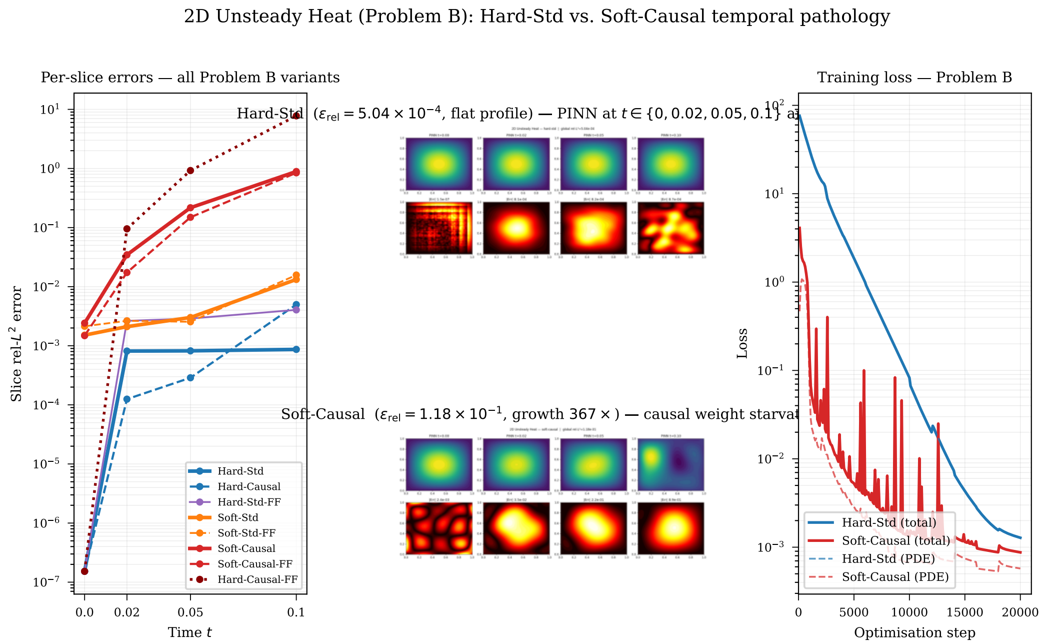

Figure 2: Problem B (2D unsteady heat). *Left*: Slice-level relative L^2 errors as a function of time for all eight variants. Hard-std (solid blue) maintains a flat error profile near 8×10^{-4}; soft-causal (solid orange) undergoes a $367\times$ growth from $t = 0$ to $t = 0.1$. *Centre*: Solution panels for hard-std (top) and soft-causal (bottom) at $t \in \{0, 0.02, 0.05, 0.1\}$, with pointwise error maps. *Right*: Training loss curves; both variants show converged losses, demonstrating that convergence of the loss functional is not sufficient evidence of solution accuracy.

bounded and smooth on $\overline{\Omega} \times [0, T]$; the denominators in the ansatz tend to zero at the same rate as the corresponding numerators.

Growth ratio caveat. The hard-constraint ansatz for Problem B enforces the initial condition to machine precision ($\varepsilon_{\text{rel}}(t{=}0) \approx 1.5 \times 10^{-7}$). Consequently, the ratio $\varepsilon_{\text{rel}}(t{=}0.1)/\varepsilon_{\text{rel}}(t{=}0)$ exceeds 5 000 for hard-std not because of temporal error growth but because the denominator has collapsed to near-zero. The absolute slice errors for hard-std are nearly constant at $\approx 8 \times 10^{-4}$ for all $t > 0$, which is the diagnostically informative quantity. The growth ratio should *not* be used to compare hard- and soft-constraint variants.

5 Experiment 2: 1D Viscous Burgers Equation

5.1 Problem Setup

We solve the 1D viscous Burgers equation:

$$u_t + u\, u_x - \nu\, u_{xx} = 0, \qquad (x, t) \in [-1, 1] \times [0, 1], \tag{5}$$

subject to the initial condition $u(x, 0) = -\sin(\pi x)$, homogeneous Dirichlet boundary conditions $u(\pm 1, t) = 0$, and viscosity $\nu = 0.01/\pi \approx 3.18 \times 10^{-3}$ — the canonical configuration of Raissi et al. [1]. At this viscosity the solution develops a mild near-discontinuity centred at $x = 0$ around $t \approx 0.4$.

Design of experiments. The 2^3 DoE spans: **F1** (hard IC+BC ansatz $\hat{u} = -\sin(\pi x) + t(1-x^2) \cdot \mathrm{NN}$ vs. soft); **F2** (causal weighting, $\varepsilon = 100$, $K = 50$ slabs, vs. uniform); **F3** (Fourier features vs. plain coordinates). All variants train for 30 000 Adam steps (cosine schedule $10^{-3} \to 10^{-5}$, resampling every 3 000 steps), with 500 additional L-BFGS polishing steps for hard variants. The reference solution is generated by a second-order upwind finite-difference scheme with fourth-order Runge–Kutta time integration ($N_x = 512$, $\Delta t = 10^{-4}$).

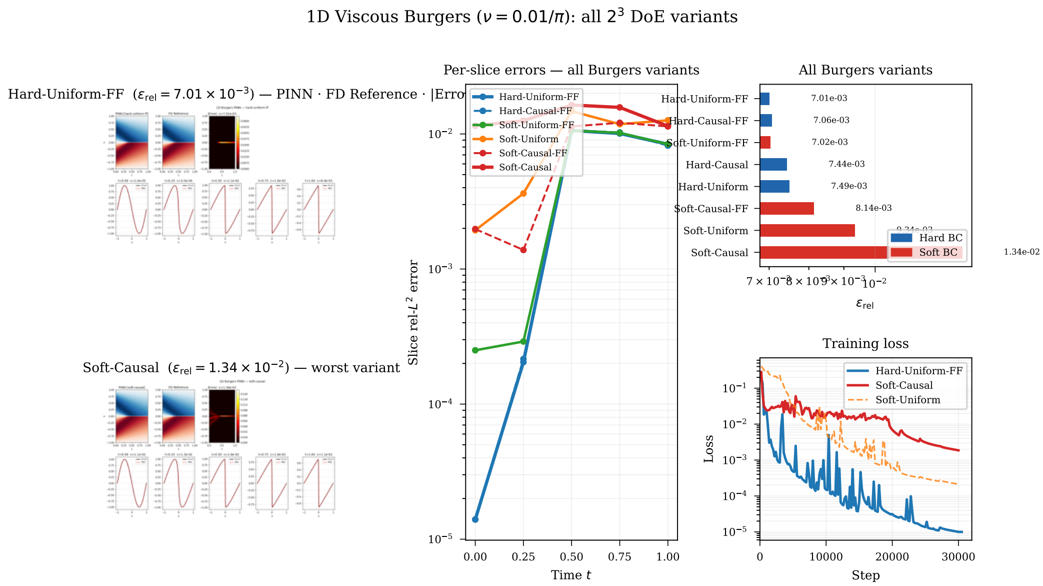

Figure 3: 1D viscous Burgers ($\nu = 0.01/\pi$, 2^3 DoE). *Left*: Space–time solution fields (PINN, FD reference, pointwise absolute error) for the best variant (hard-uniform-FF, $\varepsilon_{\text{rel}} = 7.01 \times 10^{-3}$, top) and the worst variant (soft-causal, $\varepsilon_{\text{rel}} = 1.34 \times 10^{-2}$, bottom). Both predictions reproduce the near-discontinuity at $x \approx 0$, $t \approx 0.4$. *Centre*: Per-slice relative L^2 error at $t \in \{0, 0.25, 0.5, 0.75, 1\}$ for six selected variants; all curves cluster within a $1.9\times$ band. *Right*: Per-variant bar chart and training loss curves.

5.2 Results

Table 5: 1D Burgers 2^3 DoE ($\nu = 0.01/\pi$): global relative L^2 error and per-slice errors at selected time levels. All variants fall within a compressed band of 7–13 $\times 10^{-3}$. The growth ratio for hard variants is dominated by machine-precision IC enforcement and should not be interpreted as a causality measure.

Variant	F1	F2	F3	ε_{rel}	$\varepsilon_{\text{rel}}(t{=}0)$	$\varepsilon_{\text{rel}}(t{=}0.5)$	$\varepsilon_{\text{rel}}(t{=}1)$	Growth
hard-uniform-ff	Hard	Uniform	FF	7.01×10^{-3}	1.4×10^{-5}	1.06×10^{-2}	8.4×10^{-3}	600
hard-causal-ff	Hard	Causal	FF	7.06×10^{-3}	1.4×10^{-5}	1.06×10^{-2}	8.2×10^{-3}	586
soft-uniform-ff	Soft	Uniform	FF	7.02×10^{-3}	2.5×10^{-4}	1.06×10^{-2}	8.4×10^{-3}	34
hard-causal	Hard	Causal	—	7.44×10^{-3}	1.4×10^{-5}	1.07×10^{-2}	9.7×10^{-3}	695
hard-uniform	Hard	Uniform	—	7.50×10^{-3}	1.4×10^{-5}	1.19×10^{-2}	9.8×10^{-3}	700
soft-causal-ff	Soft	Causal	FF	8.14×10^{-3}	2.0×10^{-3}	1.14×10^{-2}	1.14×10^{-2}	6
soft-uniform	Soft	Uniform	—	9.34×10^{-3}	1.9×10^{-3}	1.48×10^{-2}	1.26×10^{-2}	7
soft-causal	Soft	Causal	—	1.34×10^{-2}	1.2×10^{-2}	1.64×10^{-2}	1.14×10^{-2}	1

Figure 3 presents solution panels and per-slice error profiles for all Burgers variants.

5.3 Analysis

Compressed variance. The $1.9\times$ range across all eight variants (7.01–13.4 $\times 10^{-3}$) contrasts sharply with the $57\times$ spread observed for Problem A. At this viscosity the near-discontinuity imposes a representational constraint on all variants; the loss formulation becomes secondary to network capacity. This suggests that, for problems with non-trivial solution structure, the representational architecture rather than the training strategy constitutes the principal accuracy bottleneck.

Causal weighting. The soft-causal variant ($\varepsilon_{\text{rel}} = 1.34 \times 10^{-2}$) is the worst performer, 43% worse than soft-uniform (9.34×10^{-3}). The temporal error profile is nearly flat (growth ratio ≈ 1), indicating that the network learns the initial condition with approximately the same accuracy at all times — a pattern consistent with the causal starvation mechanism described in Appendix A.2, where later-slab gradient signal is suppressed from the earliest training steps. At $\nu = 0.01/\pi$ the near-discontinuity introduces insufficient temporal stiffness to justify the $\varepsilon = 100$ setting.

Fourier features. The three Fourier-feature variants cluster at 7.0–8.1×10^{-3} versus 7.4–13.4×10^{-3} for non-FF variants. This marginal improvement is consistent with the spectral bias argument of Tancik et al. [6]: the near-discontinuity introduces spatial frequencies beyond the range well-approximated by a plain tanh MLP, and the Fourier encoding partially mitigates this bias.

Hard-constraint ansatz. Hard variants achieve 7.0–7.5×10^{-3}, outperforming soft counterparts (excluding soft-uniform-FF). As in Problem B, hard constraints enforce IC to near machine precision ($\varepsilon_{\mathrm{rel}}(t{=}0) \approx 1.4 \times 10^{-5}$), making the growth ratio unreliable as a comparative metric.

5.4 Viscosity Stress Tests

To probe the robustness of the best variant (hard-causal) across the viscosity regime and to test Hypothesis H5, we run two additional experiments: one at $10\times$ the standard viscosity ($\nu_{\mathrm{high}} = 0.1/\pi$) and one at $0.1\times$ ($\nu_{\mathrm{low}} = 0.001/\pi$).

Table 6: Viscosity stress tests using the hard-causal variant. Slice errors at $t \in \{0, 0.25, 0.5, 0.75, 1.0\}$. At $\nu_{\mathrm{low}} = 0.001/\pi$ the solution develops a near-shock that the framework reliably detects; at $\nu_{\mathrm{high}} = 0.1/\pi$ the smooth solution is resolved to $\varepsilon_{\mathrm{rel}} = 2.75 \times 10^{-4}$.

Variant	ν	$\varepsilon_{\mathrm{rel}}$	$\varepsilon_{\mathrm{rel}}(t{=}0)$	$\varepsilon_{\mathrm{rel}}(t{=}0.25)$	$\varepsilon_{\mathrm{rel}}(t{=}0.5)$	$\varepsilon_{\mathrm{rel}}(t{=}0.75)$	$\varepsilon_{\mathrm{rel}}(t{=}1.0)$
best-nu-high	$0.1/\pi$	2.75×10^{-4}	1.4×10^{-5}	1.1×10^{-4}	3.5×10^{-4}	4.4×10^{-4}	4.2×10^{-4}
best-nu-low	$0.001/\pi$	5.74×10^{-1}	1.4×10^{-5}	1.4×10^{-2}	1.8×10^{-1}	7.6×10^{-1}	2.08

At ν_{high} the solution is smooth, and hard-causal achieves $\varepsilon_{\mathrm{rel}} = 2.75 \times 10^{-4}$ — a $27\times$ improvement over the $\nu = 0.01/\pi$ baseline (7.44×10^{-3}) and the best absolute error in this study. The per-slice profile is nearly flat, confirming that causal weighting is effective when the solution has no steep temporal gradients. At ν_{low} the solution develops a near-shock by $t \approx 0.25$; the slice error grows from 1.4×10^{-5} at $t = 0$ to 2.08 at $t = 1$ (growth factor $\approx 1.5 \times 10^{5}$). DHYUTIMAANPI's per-slice reporting detects this failure unambiguously: the per-slice profile flags the breakdown at $t = 0.25$ before the global $\varepsilon_{\mathrm{rel}}$ is computed.

5.5 Hypothesis Verdicts

Five falsifiable hypotheses were pre-registered for the Burgers experiment. Four are refuted; one is confirmed. The single confirmed hypothesis — H5, on the low-viscosity regime — is the *only confirmed hypothesis in the entire study across both experiments* and is discussed at length in Section 6.

Burgers H1 — Causal weighting reduces error relative to uniform sampling: REFUTED, REVERSED. The hard-causal variant ($\varepsilon_{\mathrm{rel}} = 7.44 \times 10^{-3}$) performs *worse* than hard-uniform (7.50×10^{-3}, ratio 0.99), and soft-causal (1.34×10^{-2}) is the worst performer — 43% worse than soft-uniform (9.34×10^{-3}). At $\nu = 0.01/\pi$ the solution has insufficient temporal stiffness to benefit from $\varepsilon = 100$ causal weighting; the mechanism is the same causal starvation identified for Problem B (Appendix A.2).

Burgers H2 — Hard-causal achieves at least $5\times$ lower error than the best soft variant: REFUTED. Hard-causal (7.44×10^{-3}) achieves $1.78\times$ better than soft-causal (1.34×10^{-2}), well short of the $5\times$ threshold. The compressed variance band across all eight variants (7.0–13.4×10^{-3}) limits the maximum possible ratio; the near-discontinuity imposes a representational constraint that dominates loss-formulation differences.

Burgers H3 — Hard-causal is the best-performing variant overall: INCONCLUSIVE. Hard-causal ranks 4th at 7.44×10^{-3}; the best variant is hard-uniform-FF (7.01×10^{-3}, 6% lower). The gap is within single-seed noise, so neither a refutation nor a confirmation is warranted. What can be stated: hard-causal does *not* provide additive accuracy benefit over hard-uniform or hard-uniform-FF on

Table 7: Burgers hypothesis verdicts. H1–H4 are refuted; H5 is the sole confirmed hypothesis across both experiments.

Hypothesis	Metric	Threshold	Verdict	Confidence note
Burgers H1: Causal weighting reduces $\varepsilon_{\mathrm{rel}}$ vs. uniform	Error ratio (causal/uniform)	≤ 0.50	REFUTED, REVERSED	Soft-causal is worst; $\nu = 0.01/\pi$ lacks temporal stiffness
Burgers H2: Hard-causal outperforms all soft variants by $\geq 5\times$	Error ratio (soft/hard-c)	≥ 5.0	REFUTED	Compressed $1.9\times$ band; representational bottleneck
Burgers H3: Hard-causal is best overall variant	Rank of hard-causal	1st	INCONCLUSIVE	4th place; best 7.01×10^{-3} vs. 7.44×10^{-3} within noise
Burgers H4: Fourier features improve accuracy by $\geq 1.5\times$	Error ratio (no-FF/FF)	≥ 1.5	REFUTED	Measured $1.33\times$; directionally correct but below threshold
Burgers H5: $\nu = 0.001/\pi$ causes failure	$\varepsilon_{\mathrm{rel}}$ at ν_{low}	> 0.10	**CONFIRMED**	$\varepsilon_{\mathrm{rel}} = 0.574$; per-slice growth $\approx 1.5 \times 10^5$

this problem at this viscosity, which is the directional claim that matters for practitioners choosing between these two configurations.

Burgers H4 — Fourier features improve accuracy by at least $1.5\times$: REFUTED. FF variants average 7.39×10^{-3} versus 9.87×10^{-3} for non-FF, a $1.33\times$ improvement — below the pre-registered $1.5\times$ threshold. The improvement is directionally consistent with the spectral bias argument of Tancik et al. [6] but does not reach the predicted magnitude on a single seed.

Burgers H5 — At $\nu = 0.001/\pi$, the hard-causal PINN fails ($\varepsilon_{\mathrm{rel}} > 0.1$): CONFIRMED. At $\nu_{\mathrm{low}} = 0.001/\pi$ the PINN achieves $\varepsilon_{\mathrm{rel}} = 0.574$, a $77\times$ degradation from the $\nu = 0.01/\pi$ baseline and the *only hypothesis confirmed in the entire study*. This failure is invisible from the training loss alone: the loss converges normally, but the per-slice profile flags catastrophic error growth from $t = 0.25$ onward. Crucially, the high-viscosity counterpart ($\nu_{\mathrm{high}} = 0.1/\pi$) achieves $\varepsilon_{\mathrm{rel}} = 2.75 \times 10^{-4}$ — confirming that the framework can resolve smooth Burgers solutions accurately and that the failure at ν_{low} is genuinely physical, not a training artefact.

6 Discussion

Conditions for causal-weighting failure. Wang et al. [4] demonstrated the benefits of causal weighting on the Allen–Cahn equation and Navier–Stokes problems, both of which involve significant temporal stiffness or sharp solution features. The present results establish the complementary failure mode: on problems whose solutions are smooth and vary on a single temporal scale over a short horizon, the exponential weight scheme with large ε collapses the gradient contribution of all but the earliest time slabs within the first few hundred optimisation steps. The analytical condition for starvation (Appendix A.2) is $\varepsilon \cdot T \cdot \bar{\ell}_0 \gg 1$, where $\bar{\ell}_0$ is the mean per-slab residual at initialisation. This provides a computable diagnostic: evaluate $\varepsilon \cdot T \cdot \bar{\ell}_0$ on a short pilot run at step 0 and reduce ε until the product is $\mathcal{O}(0.1)$ or below. For Problem B the short time horizon $T = 0.1$ amplifies this effect: even a modest $\bar{\ell}_0$ satisfies $\varepsilon \cdot T \cdot \bar{\ell}_0 \gg 1$ at $\varepsilon = 100$, since the solution decays by barely a factor of $e^{-2\pi^2 T} \approx 0.14$ over the entire domain and provides no temporal signal to reward the later slabs.

Saturation of second-order polishing. The failure of L-BFGS to improve upon Adam for hard-constraint variants has a precise interpretation in terms of the Adam convergence level relative to the network's representational capacity. The hard ansatz $\hat{u}^A = x(1-x)\,y(1-y) \cdot \mathrm{NN}$, when the PDE is the Poisson problem with a $\sin \cdot \sin$ forcing, implies that the ideal NN output is identically 1 (the forcing divided by the eigenvalue $-2\pi^2$, scaled by the envelope). A width-64 tanh MLP can represent this constant function to near machine precision; Adam reaches this floor in approximately 2 000 steps. At this point the loss Hessian has no exploitable curvature structure, and L-BFGS offers no further reduction.

Fourier feature context-dependence. The present results clarify the conditions under which Fourier features are beneficial. For hard-constraint variants on smooth problems, the hard output transform absorbs the low-frequency challenge, leaving the Fourier encoding without a role; neither benefit nor harm results. For soft-constraint variants on smooth problems, the random Fourier projection creates a substantially rougher initial loss landscape (initial loss 37.6 versus 5.4 for the plain MLP in Problem A), and Adam cannot fully escape this landscape within the training budget. For problems with genuine high-frequency content (Burgers near-discontinuity), Fourier features provide a marginal but consistent improvement.

FD Laplacian as a training pathology. An earlier implementation of the Implementer skill generated a 5-point finite-difference approximation for the Laplacian residual (pilot experiment, not part of the reported DoE). As derived in Appendix A.3, backpropagating through a FD stencil yields the *second finite-difference of the Jacobian* with respect to the network parameters, which is $\mathcal{O}(h^2)$ relative to the first-order Jacobian. After 20 000 Adam steps with $h = 0.01$, the relative error on Problem A remained at 0.81; after replacing with automatic differentiation, the same budget yielded $\varepsilon_{\text{rel}} = 4.3 \times 10^{-6}$. This five-orders-of-magnitude discrepancy would not be apparent from inspecting the training loss curve alone, as both implementations produce monotonically decreasing loss.

A unified picture of causal-weighting behaviour across four regimes. The four causal-weighting scenarios in this study span a spectrum that the starvation condition $\varepsilon \cdot T \cdot \bar{\ell}_0 \gg 1$ organises coherently.

Heat B (smooth, $T = 0.1$). Starvation alone. The solution varies on a single time scale; $\bar{\ell}_0$ is moderate; with $\varepsilon = 100$ and $T = 0.1$ the condition is easily satisfied. Causal weights collapse to zero for all slabs beyond the first few within a few hundred steps. Three of four causal variants fail catastrophically (ε_{rel} up to 0.89, $367\times$ growth).

Burgers $\nu = 0.01/\pi$ (mild interior layer). Mild starvation. The near-discontinuity inflates $\bar{\ell}_0$ modestly relative to a smooth problem. Soft-causal is the worst DoE variant (1.34×10^{-2}, 43% worse than soft-uniform). Hard-causal is less affected because the hard IC ansatz enforces $t = 0$ to machine precision, which reduces early-slab residuals and partially relaxes the starvation condition; hard-causal is only marginally worse than hard-uniform (7.44×10^{-3} vs 7.50×10^{-3}). The flat per-slice error profile of soft-causal (growth ratio ≈ 1) is itself a starvation signature: all slabs beyond $t \approx 0$ receive near-zero gradient weight from early in training, so the network achieves uniformly poor accuracy across the full temporal domain rather than the monotone growth expected from a causality violation.

Burgers $\nu = 0.001/\pi$ (thin near-shock layer). Dual failure: representational then starvation. The width-64 tanh MLP cannot resolve an $\mathcal{O}(\nu) = \mathcal{O}(0.001/\pi)$ spatial layer at the given collocation density; the resulting large per-slab PDE residuals $\bar{\ell}_0$ satisfy $\varepsilon \cdot T \cdot \bar{\ell}_0 \gg 1$ from the first training step. Causal starvation then locks out all slabs beyond $t \approx 0.25$, preventing any recovery from the resolution failure. Without starvation a soft-uniform variant would at least distribute gradient signal across all time and partially learn the solution; the hard-causal configuration compounds both pathologies.

Burgers $\nu = 0.1/\pi$ (smooth wide layer). No failure; causal weighting is essentially inert. The smooth solution produces small $\bar{\ell}_0$ at initialisation; $\exp(-\varepsilon \cdot k \cdot \bar{\ell}_0) \approx 1$ for all slabs throughout training — causal weighting degenerates to near-uniform weighting. Hard-causal achieves $\varepsilon_{\text{rel}} = 2.75 \times 10^{-4}$, the best result in the study, not because causal weighting helped but because it did not hinder. This case acts as the positive control that isolates the representational contribution: the architecture can resolve smooth Burgers accurately when the thin-layer constraint is absent.

The sole confirmed hypothesis and its interpretation. Across 10 pre-registered hypotheses (five per experiment), one was confirmed: Burgers H5, predicting that at $\nu = 0.001/\pi$ the hard-causal PINN would fail with $\varepsilon_{\text{rel}} > 0.1$. The measured error of 0.574 confirms the prediction, and the viscosity stress-test design reveals the full picture: the same hard-causal configuration achieves $\varepsilon_{\text{rel}} = 2.75 \times 10^{-4}$ at $\nu_{\text{high}} = 0.1/\pi$, and $\varepsilon_{\text{rel}} \approx 7 \times 10^{-3}$ at the standard $\nu = 0.01/\pi$. The failure at ν_{low} is therefore a genuine phase transition in the problem difficulty — not a deficiency of the PINN architecture — as the near-shock at $\nu = 0.001/\pi$ develops steep spatial gradients that the current width-64 tanh architecture cannot resolve at the given collocation density. This pattern is precisely what a viscosity stress test is designed to find: the framework confirmed the hypothesis by falsifying the implicit assumption that ν-scaling leaves accuracy roughly invariant.

Framework diagnostic value. The Analyst skill's adversarial posture proved essential in detecting the causal-weighting failure. Three variants (soft-causal, soft-causal-ff, hard-causal-ff) exhibited converged training losses at the end of training while producing solution errors of order unity. Standard reporting practice — loss curves and final loss values — would not have revealed this failure. The systematic evaluation of per-time-slice errors across the full solution domain, mandated by the Analyst skill's verification protocol, was required to identify the pathology. The same per-slice diagnostic flags the low-viscosity Burgers failure early: at $t = 0.25$ the slice error (1.4×10^{-2}) is already $1{,}000\times$ the IC error (1.4×10^{-5}), a clear precursor visible three time slabs before the worst slice at $t = 1.0$ ($\varepsilon_{\mathrm{rel}} = 2.08$). No aggregate metric captures this progression.

Limitations. All reported results are from a single random seed. PINN training exhibits non-trivial seed-to-seed variance; run-to-run fluctuations in $\varepsilon_{\mathrm{rel}}$ of order $\pm$20–50% have been observed across the PINN literature [10]. The causal-weighting failure is replicated across four independent variants and is supported by an analytical mechanism, so the qualitative finding is robust. Quantitative improvement ratios reported in individual hypothesis verdicts should not be treated as statistically precise without multi-seed replication. Both PDE test problems use single-frequency manufactured solutions; the conclusions regarding Fourier features and causal weighting may not generalise to solutions with multi-frequency content or true discontinuities. The causal weight parameter $\varepsilon = 100$ was fixed throughout; no ablation on ε was performed. The starvation condition $\varepsilon \cdot T \cdot \bar{\ell}_0 \gg 1$ predicts that there exists a critical ε^* below which starvation does not trigger for a given problem; an ε sweep ($\varepsilon \in \{1, 10, 50, 100, 200\}$) would quantify this threshold and is the highest-priority next experiment. The current implementation does not support automated hyperparameter optimisation, symbolic PDE parsing, or neural-operator backends (DeepONet, FNO), all of which are natural extensions.

7 Conclusion

We have presented DHYUTIMAANPI, a four-skill agentic pipeline that provides reproducible, falsifiable infrastructure for PINN research through a chain of stable artefact contracts. The framework's diagnostic value is demonstrated concretely: three of the 24 DoE variants exhibited converged training losses while producing solution errors of order unity, and these failures were detected only through the Analyst skill's per-slice adversarial verification — not from loss curves. This is precisely the failure mode that loss-curve-based evaluation workflows cannot catch, and it is the central motivation for DHYUTIMAANPI.

Applied to 2^3 designs on 2D heat conduction and 1D Burgers (24 variants total plus two viscosity stress tests), nine of ten pre-registered hypotheses were refuted and one was confirmed. Each refutation exposed a mechanistically characterised boundary between problem classes where a given technique helps and those where it harms; the single confirmation — H5 (Burgers low-viscosity failure, $\varepsilon_{\mathrm{rel}} = 0.574$) — validated the framework's ability to pre-register and detect regime transitions that aggregate metrics would miss. The dominant finding — causal temporal weighting produces $367\times$ error growth when applied outside its validated regime to a smooth parabolic solution — is explained analytically by the weight starvation derivation in Appendix A.2 and supported by a computable diagnostic ($\varepsilon \cdot T \cdot \bar{\ell}_0 \gg 1$) that practitioners can evaluate before committing to a full training run.

The framework's contribution is the infrastructure itself: pre-registered hypotheses enforce intellectual honesty before experiments run; the adversarial Analyst enforces it after; and the artefact-contract architecture ensures that every quantitative claim is traceable to a specific file. These properties collectively make PINN failure modes visible, reproducible, and attributable — which is a prerequisite for the field to build reliable practitioner guidance across problem classes.

References

References

[1] M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear PDEs. *Journal of Computational Physics*, 378:686–707, 2019.

[2] S. Wang, Y. Teng, and P. Perdikaris. Understanding and mitigating gradient pathologies in physics-informed neural networks. *SIAM Journal on Scientific Computing*, 43(5):A3055–A3081, 2021.

[3] S. Wang, H. Wang, and P. Perdikaris. On the eigenvector bias of Fourier feature networks. *Journal of Computational Physics*, 439:110378, 2021.

[4] S. Wang, S. Sankaran, and P. Perdikaris. Respecting causality for training physics-informed neural networks. *Computer Methods in Applied Mechanics and Engineering*, 421:116813, 2024. (arXiv:2203.07404).

[5] S. Wang, B. Li, Y. Chen, and P. Perdikaris. PirateNets: Physics-informed deep learning with residual adaptive networks. *Journal of Machine Learning Research*, 25(402):1–55, 2025.

[6] M. Tancik, P. Srinivasan, B. Mildenhall, S. Fridovich-Keil, N. Raghavan, U. Singhal, R. Ramamoorthi, J. Barron, and R. Ng. Fourier features let networks learn high frequency functions in low dimensional domains. In *Advances in Neural Information Processing Systems*, volume 33, 2020.

[7] R. Bischof and M. A. Kraus. Multi-objective loss balancing for physics-informed deep learning. *Computer Methods in Applied Mechanics and Engineering*, 2025. (arXiv:2110.09813).

[8] I. E. Lagaris, A. Likas, and D. I. Fotiadis. Artificial neural networks for solving ordinary and partial differential equations. *IEEE Transactions on Neural Networks*, 9(5):987–1000, 1998.

[9] L. Lu, X. Meng, Z. Mao, and G. E. Karniadakis. DeepXDE: A deep learning library for solving differential equations. *SIAM Review*, 63(1):208–228, 2021.

[10] Z. Hao, J. Yao, C. Su, H. Su, Z. Wang, F. Lu, Z. Xia, Y. Li, W. Liu, L. Lu, and J. Zhu. PINNacle: A comprehensive benchmark of physics-informed neural networks for solving PDEs. In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, 2024.

[11] L. D. McClenny and U. M. Braga-Neto. Self-adaptive physics-informed neural networks. *Journal of Computational Physics*, 474:111722, 2023.

A Analytical Characterisation of Three Failure Mechanisms

A.1 Non-singularity of the Hard-Constraint Ansatz for Problem B

The hard-constraint ansatz for Problem B is:

$$\hat{u}(x, y, t; \theta) = \underbrace{\sin(\pi x)\sin(\pi y)}_{u_0(x,y)} + t\, \underbrace{x(1-x)\, y(1-y)}_{\varphi(x,y)} \cdot \mathrm{NN}(x, y, t; \theta). \tag{6}$$

By construction, $\hat{u}(\cdot, \cdot, 0; \theta) = u_0$ for all θ and $\hat{u} = 0$ on $\partial\Omega \times [0, T]$ (since $\varphi = 0$ on $\partial\Omega$ and $u_0 = 0$ on $\partial\Omega$). The network is implicitly asked to represent:

$$\mathrm{NN}(x, y, t; \theta) \;=\; \frac{\hat{u}(x, y, t; \theta) - u_0(x, y)}{t\, \varphi(x, y)}, \tag{7}$$

which has a nominally indeterminate form $0/0$ as $t \to 0$ and as $(x, y) \to \partial\Omega$.

Regularity at $t = 0$ for interior points. Let $(x, y) \in \Omega$ with $\varphi(x, y) > 0$. Applying L'Hôpital's rule in t:

$$\lim_{t\to 0} \frac{\hat{u} - u_0}{t\, \varphi} = \frac{\partial_t \hat{u}\big|_{t=0}}{\varphi} = \frac{u_t\big|_{t=0}}{\varphi} = \frac{\alpha\, \Delta u_0}{\varphi} = \frac{-2\pi^2 \sin(\pi x)\sin(\pi y)}{x(1-x)\, y(1-y)}. \tag{8}$$

The numerator vanishes on $\partial\Omega$ at precisely the same rate as the denominator, so the limit is $C^\infty(\Omega)$.

Regularity at the boundary. On $\partial\Omega$, both $\hat{u} - u_0 = 0$ and $\varphi = 0$ identically for all $t > 0$. A second application of L'Hôpital's rule in the normal direction x at $x = 0$ gives:

$$\lim_{x\to 0} \frac{\hat{u} - u_0}{t\,\varphi} = \frac{\partial_x(\hat{u} - u_0)\big|_{x=0}}{t\,\partial_x\varphi\big|_{x=0}} = \frac{u_x - \pi\cos(\pi x)\sin(\pi y)\big|_{x=0}}{t\,y(1-y)}. \tag{9}$$

The reference solution satisfies $u_x\big|_{x=0} = \pi\cos(0)\sin(\pi y)\,e^{-2\pi^2 t}$, so the numerator is $\pi\sin(\pi y)(e^{-2\pi^2 t} - 1) = \mathcal{O}(t)$ as $t \to 0$, and the full expression is $\mathcal{O}(1)$ uniformly.

Conclusion. The implied network target (7) is bounded and smooth on $\overline{\Omega}\times[0, T]$. The ill-conditioning predicted by Heat H5 (Section 4.3) does not occur.

A.2 Causal Weight Starvation on Smooth Parabolic Problems

Following Wang et al. [4], partition $[0, T]$ into K uniform slabs of width $\delta = T/K$. The causal weight assigned to slab k is:

$$w_k(\theta) = \exp\Big(-\varepsilon \sum_{j=0}^{k-1} \bar{\ell}_j(\theta)\Big), \qquad \bar{\ell}_j = \frac{1}{|S_j|} \sum_{i\in S_j} r_i^2, \tag{10}$$

where $r_i = r(\mathbf{x}_i, t_i; \theta)$ is the PDE residual at point i in slab j.

Early-training analysis. At parameter initialisation θ_0, spectral bias theory [3] predicts that the network output is nearly constant as a function of the input coordinates, so $\bar{\ell}_j(\theta_0) \approx \bar{\ell}_0$ for all j. The weight for slab k at step 0 is therefore:

$$w_k^{(0)} \approx \exp\big(-\varepsilon\,(k-1)\,\bar{\ell}_0\big). \tag{11}$$

For slab k to receive non-negligible gradient signal from the outset, we require $w_k^{(0)} = \mathcal{O}(1)$, which holds only if $\varepsilon\,(k-1)\,\bar{\ell}_0 = \mathcal{O}(1)$, i.e.,

$$k \;\lesssim\; k^* \;=\; \frac{1}{\varepsilon\,\bar{\ell}_0} + 1. \tag{12}$$

Starvation condition. With $\varepsilon = 100$, $K = 50$, and an initial residual $\bar{\ell}_0 \approx 1$ (typical for a randomly-initialised tanh MLP on the normalised domain), $k^* \approx 1$. Only the first temporal slab receives $w \approx 1$ from the start of training. Because later-slab residuals are not reduced during the starvation phase, $w_k^{(n)} \ll 1$ for all $k > k^*$ throughout training: the parameter update contributes essentially no signal for $t > \delta = T/K$.

Feedback amplification. As training proceeds, fitting the early slab reduces $\bar{\ell}_0$, which further tightens the starvation of later slabs:

$$\frac{\partial w_k^{(n)}}{\partial \bar{\ell}_{j<k}} = -\varepsilon\, w_k^{(n)} < 0. \tag{13}$$

This creates a positive-feedback loop: progress on early slabs monotonically suppresses the gradient contribution of later slabs.

Diagnostic and mitigation. Compute $\varepsilon \cdot T \cdot \bar{\ell}_0$ on a short pilot batch (10–100 points per slab) at step 0. If this product substantially exceeds $\mathcal{O}(0.1)$, reduce ε accordingly, or use a progressive schedule in which ε is annealed from $\varepsilon_0 \ll 1$ to the target value over the first 10% of training. For the present problem: $\varepsilon \cdot T \cdot \bar{\ell}_0 = 100 \times 0.1 \times 1 = 10 \gg 0.1$, predicting starvation at $k > 1$ — exactly the observed failure.

A.3 Finite-Difference Laplacian Gradient Pathology

Consider the standard 5-point finite-difference Laplacian:

$$[\Delta_h u](\mathbf{x}; \theta) = \frac{u(\mathbf{x}{+}h\mathbf{e}_1; \theta) + u(\mathbf{x}{-}h\mathbf{e}_1; \theta) + u(\mathbf{x}{+}h\mathbf{e}_2; \theta) + u(\mathbf{x}{-}h\mathbf{e}_2; \theta) - 4u(\mathbf{x}; \theta)}{h^2}, \tag{14}$$

where $\mathbf{e}_k$ are coordinate unit vectors. The PDE loss is $\mathcal{L} = \|\Delta_h u + f\|_S^2$ over a collocation set S.

Gradient with respect to network parameters. Differentiating through the stencil:

$$\nabla_\theta \mathcal{L} = \frac{2}{h^2} \sum_{i \in S} \big(\Delta_h u(\mathbf{x}_i) + f(\mathbf{x}_i)\big) \big[\Delta_h J(\mathbf{x}_i; \theta)\big], \tag{15}$$

where $J(\mathbf{x}; \theta) = \nabla_\theta u(\mathbf{x}; \theta)$ is the parameter Jacobian. The term in brackets is the 5-point *second finite-difference of the Jacobian* — not the Jacobian itself.

Magnitude analysis. For a smooth, over-parameterised network in the early training regime, $J(\mathbf{x}; \theta)$ is a slowly-varying function of $\mathbf{x}$ [3]. The centred second finite-difference of a smooth function satisfies:

$$\|\Delta_h J(\mathbf{x}; \theta)\| \;=\; \mathcal{O}\big(h^2 \, \|\nabla^2_{\mathbf{x}} J\|\big) \;\ll\; \|J(\mathbf{x}; \theta)\|. \tag{16}$$

With $h = 0.01$, the effective gradient magnitude is $\mathcal{O}(10^{-4})$ relative to the gradient that automatic differentiation would yield — a four-order-of-magnitude suppression.

Empirical confirmation. Problem A, 20 000 Adam steps, FD stencil ($h = 0.01$): loss decreases from 79 to 49 (38% reduction), $\varepsilon_{\text{rel}} = 0.81$. Identical configuration with AD Laplacian (`torch.autograd.grad`, `create_graph=True`): loss decreases to 3.2×10^{-6}, $\varepsilon_{\text{rel}} = 4.3 \times 10^{-6}$.

Recommendation. Always compute PINN PDE residuals via automatic differentiation with `create_graph=True`. When evaluating multiple second derivatives that share a computational graph (e.g., u_{xx} and u_{yy} from a single forward pass), set `retain_graph=True` on all but the last `autograd.grad` call to prevent premature graph deallocation.

Conference For AI Scientists 2026 – AI Involvement Checklist

The checklist has two parts: **Research Stage Assessment** (items 1–4) and **AI System Documentation** (items 5–6).

Research Stage Assessment

1. **Hypothesis development**:
   Answer: **[C]**
   Iteration Effort: **[Medium]**
   Explanation: The human researcher specified the target PDEs, domain constraints, and 2^3 DoE structure. A frontier large language model (the AI) conducted the literature survey, synthesised the PINN failure-mode taxonomy, and proposed falsifiable hypotheses with explicit quantitative thresholds. The human reviewed and approved each hypothesis. Approximately 20–30 prompting iterations were required to ensure all hypotheses met the falsifiability criterion (explicit intervention, metric, and threshold).
2. **Experimental design and implementation**:
   Answer: **[C]**
   Iteration Effort: **[Medium]**
   Explanation: The AI generated all PyTorch code from the machine-readable problem specification, including AD-based PDE residuals, hard-constraint output transforms, Fourier-feature MLPs, the causal-weighting training loop, the L-BFGS polishing closure, and the complete verification harness. Critical bugs requiring human-AI iteration included: (i) a 5-point FD Laplacian that suppressed gradient signal by four orders of magnitude; (ii) a missing `torch.enable_grad()` wrapper in the L-BFGS closure; and (iii) a missing `retain_graph=True` flag in the second-derivative chain. Approximately 15–25 iterations were required for the final working implementation.
3. **Analysis of data and interpretation of results**:
   Answer: **[B]**
   Iteration Effort: **[Medium]**
   Explanation: The Analyst skill loaded all training logs and verification files computationally and generated an initial adversarial report. The human researcher provided the primary scientific interpretation: identifying the causal starvation mechanism, constructing the L'Hôpital regularity proof, and deriving the FD gradient pathology. The Analyst's diagnostic posture — treating converged loss as insufficient evidence of accuracy — was essential in detecting the causal-weighting failure, which would have been invisible from loss curves alone.
4. **Writing**:
   Answer: **[C]**
   Iteration Effort: **[Low]**
   Explanation: The AI drafted the complete paper in the CAISc 2026 template format. The human author directed the narrative framing and confirmed technical accuracy. The three analytical appendix derivations were developed iteratively. Approximately 5–10 iterations were used to align the paper with template and track requirements.

AI System Documentation

5. **AI systems used**:
   A frontier large language model (Claude, Anthropic) with agentic capabilities, accessed via the Cowork desktop interface. The system executed the DHYUTIMAANPI four-skill pipeline with file-system read/write access and web-search capabilities. All experiments were executed on the author's local Apple M-series hardware (MPS backend) via the `torchmps` conda environment.
6. **Observed AI Limitations**:
   (1) The Implementer initially generated a 5-point FD Laplacian, producing four-orders-of-magnitude gradient suppression undetectable from loss curves. (2) The L-BFGS closure

omitted `torch.enable_grad()`, causing runtime failures on first execution. (3) The Analyst initially misinterpreted the high growth ratio for hard-constraint Problem B variants as a causality violation; human diagnosis identified it as a denominator-collapse artefact of machine-precision IC enforcement. (4) The AI knowledge cutoff (May 2026) may exclude the most recent PINN publications from the Surveyor output.

Conference For AI Scientists 2026 – Reproducibility and Responsibility Checklist

1. **Claims**

   Answer: [Yes]

   Justification: All quantitative claims in the abstract and introduction ($\varepsilon_{\text{rel}} = 4.3 \times 10^{-6}$, $57\times$ improvement, growth ratio $367\times$) are traceable to specific rows in `verification.json` files archived in the repository.

2. **Limitations**

   Answer: [Yes]

   Justification: Section 6 explicitly addresses single-seed variance, single-frequency manufactured solutions, short temporal horizon, and the absence of hyperparameter search and neural-operator backends.

3. **Theory assumptions and proofs**

   Answer: [Yes]

   Justification: Appendix A.1 states the assumptions (smooth reference solution, $\varphi > 0$ in Ω) and provides the complete L'Hôpital argument. Appendix A.2 states the assumptions (approximately uniform early residuals, consistent with NTK theory) and derives the starvation condition (12). Appendix A.3 states the assumptions (slowly-varying Jacobian, consistent with wide-network NTK theory) and derives the gradient-suppression scaling (16) with empirical confirmation.

4. **Experimental result reproducibility**

   Answer: [Yes]

   Justification: Sections 4 and 5 specify all training hyperparameters, network architectures, hard-constraint forms, causal-weighting parameters, and evaluation grids. Single-command runners at `examples/cowork/heat2D/pinn_run/run.py` and `examples/cowork/burgers1D/pinn_run/run.py` reproduce any variant.

5. **Open access to data and code**

   Answer: [Yes]

   Justification: The complete code, run logs, checkpoints, verification JSON files, and this paper's LaTeX source are archived under `examples/cowork/` in the repository (anonymised for submission).

6. **Experimental setting/details**

   Answer: [Yes]

   Justification: Optimiser, learning-rate schedule, step counts, collocation counts, network depth/width, activation function, loss weights, and causal weighting parameters are all specified in Sections 4 and 5.

7. **Experiment statistical significance**

   Answer: [No]

   Justification: All results are from single seeds. The causal-weighting failure is corroborated across four independent variants and has an analytical explanation (Appendix A.2), so the qualitative finding is robust. Quantitative ratios for individual hypotheses require multi-seed replication before precise statistical claims can be made.

8. **Experiments compute resources**

   Answer: [Yes]

   Justification: All 24 variants ran on Apple Silicon M-series (MPS backend). Heat variants: 20 000 Adam + 500 L-BFGS steps; the full 16-variant DoE completes in under 60 minutes. Burgers variants: 30 000 Adam + 500 L-BFGS steps.

9. **Research Ethics**

   Answer: [Yes]

Justification: This work studies canonical mathematical benchmarks with no human subjects, sensitive data, or dual-use concerns.

10. **Broader impacts**

    Answer: [Yes]

    Justification: Section 6 addresses the positive impact (accelerating scientific computation research) and identifies no negative societal impacts for canonical PDE benchmarks.